

Sharing Conflict Weights Without Partitioning: A Concurrent dom/wdeg Portfolio for Word-Level Crossword Filling

[Author Name(s)]

[Department], [University]

Submitted to IEEE MIT URTC 2026

Abstract—Parallel constraint solvers generally take one of two paths: partition the search space across workers, or run an undivided portfolio of independent restarts. Prior work that shares information between constraint-satisfaction-problem (CSP) workers treats that sharing as a prelude to partitioning, gathered once and then spent. We describe a different point in this design space: a restart portfolio whose workers never partition the problem, but continuously and concurrently update a shared pool of dom/wdeg-style conflict weights throughout the entire search, with no manager and no synchronization phase. We implement this in `xfill`, a word-level CSP solver for crossword filling, and evaluate it against four other solvers – `orca-solver` (partition-based), `ingrid_core` (single-threaded reference), and two simpler, unmodified open-source backtrackers used as naive baselines – on a reproducible, fixed testbench of real and size-graded grids, plus an adversarially constructed grid designed to be exhaustively hard. Restart portfolios and partitioned search are not uniformly ordered: on real 15x15 grids, `xfill` and `orca-solver` succeed on different, only partially overlapping subsets, and oversubscribing thread count past physical core count consistently helps `xfill`, consistent with restart-based search’s known heavy-tailed runtime distribution. Shared conflict weights give a further, consistent improvement, but only once scoped to avoid two measured regressions that a naive implementation introduces. We also report a negative result: five independently plausible letter-level branching heuristics produced numerically identical output on our hardest benchmark grid, because a domain-size precondition meant none of them ever executed – caught only by direct execution-path instrumentation, not by comparing benchmark scores.

I. Introduction

Crossword filling – assigning a dictionary word to every slot in a grid so every crossing letter agrees – is a natural word-level CSP and a convenient search-research testbed: instances are human-legible, difficulty is controllable directly through grid geometry, and independently developed solvers already exist to compare against. Ginsberg’s `Dr.Fill` [1] established crossword filling as a serious weighted-CSP application over a decade ago.

This paper is not primarily about crossword filling. It uses crossword filling as the setting for a narrower question in parallel constraint search: when multiple workers run concurrently, is there value in sharing soft heuristic statistics – information that biases search order without asserting anything logically – continuously, without ever partitioning the problem those workers are solving? We

show this combination is distinct from the closest points in the literature, implement it, and measure its effect against both sophisticated and naive alternatives on a reproducible testbench.

Contributions: (1) a concurrent, decentralized conflict-weight-sharing mechanism for restart-portfolio CSP search, positioned precisely against the closest prior work; (2) a reproducible, five-solver benchmark testbench, checked into the project repository, comparing `xfill` against `orca-solver`, `ingrid_core`, and two independent open-source baselines on a fixed, size-graded grid set; and (3) a methodological negative result about validating that a heuristic under test actually executes before trusting a benchmark comparison of it.

II. Background and Related Work

A CSP assigns values to variables from finite domains subject to constraints; systematic backtracking assigns one variable at a time, propagating consistency and backtracking on a wipeout. The dom/wdeg heuristic [4] selects the variable with smallest domain-size-to-weighted-degree ratio, where a constraint’s weight increases on every wipeout it causes. Randomized backtracking exhibits heavy-tailed runtime [2], [3]: restarting periodically eliminates most of the tail and is standard in SAT and CSP solvers alike.

Two dominant parallel strategies exist. Portfolio approaches run independent searches over the whole problem concurrently; SAT portfolios such as `ManySAT` [6] additionally share learned clauses between workers without partitioning. Partitioning approaches divide the search space so each worker explores a disjoint region, giving non-overlapping coverage useful for proving unsatisfiability.

Three groups of prior CSP work sit near, but not on, the point in this design space we describe. Grimes and Wallace [5] accumulate dom/wdeg-style weights from short restarted probes of a single sequential solver, before the real search begins – never concurrent. Ehlers and Stuckey [7] parallelize a lazy-clause-generation solver by sharing nogoods and trail state between workers – concurrent, but sharing hard, logically derived facts, and requiring clause-learning machinery. Yun and Epstein’s `ELF` and `SPREAD` [8], [9] are the closest: a manager races workers under a plain dom/wdeg solver, then averages their reported

weights once each exhausts a budget, and uses the average to choose how to partition the problem for a subsequent phase. The authors state this directly: “the primary purpose of SPREAD’s portfolio phase is to glean information to support search space splitting, not to solve P ” [9], and separately that “most portfolio-based methods for CSPs do not share information” [9] at all outside their framework. A comprehensive 2018 survey [11] documents no CSP work sharing heuristic statistics continuously among concurrent, non-partitioning workers; its own taxonomy has no category for it. Embarrassingly parallel search [10], a widely used later technique, decomposes into disjoint subproblems with “almost no communication” at all – a different point again.

The mechanism in Section III occupies the gap these leave: concurrent (unlike Grimes and Wallace), soft (unlike Ehlers and Stuckey), and never used to partition (unlike Yun and Epstein).

III. System Architecture

xfill parses a grid into slots (open runs ≥ 2 cells) and crossings (cell-level intersections). Its sequential core performs dom/wdeg-style branching – selecting the slot minimizing domain size over accumulated crossing weight – then propagates letter constraints via an AC-3-style queue; a domain wipeout increases the responsible crossing’s weight.

SolveParallel runs N independent restart workers, each with its own seed and private, decaying crossing-weight table, plus one worker dedicated to a single uninterrupted exhaustive search (unlimited_budget) guaranteeing eventual completion independent of restart variance. Any worker’s sound negative result cancels the others.

Concurrent conflict-weight sharing. Alongside each worker’s private table, SolveParallel maintains one array of plain atomic counters, one per crossing, shared by every restart worker (the dedicated exhaustive worker is excluded). Every worker increments a crossing’s counter on wipeout and reads every counter when scoring candidates. The shared counters are plain, order-independent atomic increments rather than a decaying scheme – a decaying normalizer depends on the exact sequence of prior updates, race-free only with a single writer, whereas concurrent writers need an order-independent update rule. The dedicated exhaustive worker neither reads nor writes the shared array: wiring it in measurably regressed the one grid in our corpus it uniquely solved, since its entire value lies in one undisturbed trajectory. Two correctness properties were required before the resulting numbers meant anything: single-threaded calls must stay byte-identical to the unshared baseline (gated on thread count > 1), and ThreadSanitizer, run against real multi-threaded solves, reported no data races – consistent with every update being independent and order-free.

IV. A Negative Result: Branching on the Wrong Thing

Prompted by a loss against orca-solver on a difficult 7×7 grid (Section V), we implemented five letter-level branching variants – two minimum-remaining-value scoring rules, a faithful reimplement of orca-solver’s own candidate scoring (read from its published source for reference), and a reversed scan order – reasoning that a smaller per-node branching factor should help on wide-open grids.

All five produced numerically identical node and back-track counts. Five independently reasonable heuristics agreeing to the exact node is a stronger signal than “the heuristics didn’t help,” and warranted checking rather than reporting a null result. A plain counter placed directly in the branch-selection code, tracking how often the large-domain branch these heuristics govern actually executed, showed it firing zero times across more than 17,000 real search nodes: propagation from the grid’s pre-filled letters narrowed every domain below the branch’s activation threshold within the first few assignments. All five heuristics were unreachable code the entire time they were being compared, which is exactly why comparing their benchmark scores could not have distinguished them. The actual bottleneck – found only after ruling out the large-domain branch – was the small-domain branch’s fixed word ordering; extending the solver’s existing restart-randomization flag to that branch is what produced the improvements in Section V. We report this as a reusable methodological point: an A/B comparison of branching heuristics is only informative once the branch under test is confirmed, by direct instrumentation, to actually execute.

V. Experiments

Solvers. xfill (this work); orca-solver, an independently developed partition-based parallel solver treated as a black box; ingrid_core, a single-threaded reference solver; and, to contrast sophisticated architectures against unmodified naive ones, two open-source backtrackers with no restarts, no learned heuristics, and no parallelism: crossword-composer (a static-constraint-order Rust back-tracker) and savin_crossword (a textbook AC-3-plus-backtracking Python implementation). All five are run through a single fixed harness checked into this project’s repository (benchmarks/urtc2026_testbench/), against a dictionary derived deterministically from the repository’s own freely available word list (never a paid one), with a fixed random seed and a uniform 30-second cap – fully reproducible by re-running run_benchmark.py.

Grid set. A size-graded curated set (5×5 through 21×21) plus a fixed-seed random sample of 15×15 grids from a real, previously scraped corpus, all already committed to the repository.

Oversubscription. Table I shows wall time on a known-hard, previously published 7×7 grid (two pre-filled letters, no blocked cells) as thread count exceeds this machine’s 14 physical cores. Every oversubscribed configuration

TABLE I

Thread count vs. wall time on a known-hard 7×7 grid (single sample; orca-solver’s single-threaded baseline is 201.7s).

Threads	Wall time (s)
1	166.7
14	63.0
28	104.6
42	85.6
56	112.4




Fig. 1. Solve success on the size-graded curated grid set, 30s cap. crossword-composer succeeds intermittently; savin_crossword (bar height 0 throughout, per-node cost dominated by dictionary size rather than grid size) never appears.

beats orca-solver’s single-threaded baseline; 42 threads ($3 \times$ physical cores) is fastest, with contention visibly eroding the gain by 56.

Scaling and naive baselines. Across all 20 grids, xfill solved 17 (1 correct UNSAT, 2 timeouts), orca-solver 16 (1 UNSAT, 3 timeouts), ingrid_core 12, crossword-composer 5, and savin_crossword 0 – it timed out on every single grid, including the trivial 5×5 . Fig. 1 shows why the two naive baselines fail for different reasons, not the same one: crossword-composer’s failures are instance-specific rather than size-specific (it solves some 15×15 s and small grids alike, timing out unpredictably elsewhere), consistent with having no restart to recover from a single bad static variable-ordering choice computed once up front. savin_crossword’s uniform failure is not about grid size at all – its least-constraining-value step recomputes an $O(|\text{domain}|^2)$ score from scratch at every node, which our $\sim 184,000$ -word dictionary makes prohibitive regardless of grid size, since domain sizes before any grid-driven narrowing are already large. Both are real, unmodified, independently authored solvers; the gap is what restarts, adaptive heuristics, and (for xfill/orca-solver) parallelism buy at dictionary and grid sizes that matter in practice.

Real grids. Fig. 2 shows cumulative solve rate on the 15×15 sample for solvers that finish at all. Fig. 3 breaks down xfill vs. orca-solver precisely: both solve 16 of 20, xfill alone solves 1 more, both correctly prove 1 grid has no solution (not a failure for either), and 2 remain unresolved by both within the cap – the two sophisticated archi-




Fig. 2. Cumulative fraction of real 15×15 grids solved by wall time (log scale), 30s cap.




Fig. 3. Outcome overlap between xfill and orca-solver across the full grid set: the two architectures are complementary, not uniformly ordered.

tectures are complementary, not one strictly dominating, consistent with Section VI’s explanation.

Exhaustive case study. To test proof capability rather than only positive results, we constructed by hand a grid designed to be maximally constrained: five mutually crossing 15-letter entries, verified cell-by-cell against both solvers’ own parsers. Run to completion with no time limit, both solvers independently concluded no solution exists: xfill exhausted 2.16 billion nodes over 791 restarts in 76 minutes; orca-solver reported search exhausted after enumerating 224 partitions in 25 minutes – roughly $3 \times$ faster, consistent with partitioning’s non-overlapping coverage being a more direct route to a proof than restart-based search’s redundant one.

VI. Discussion

Restart portfolios benefit from heavy-tailed runtimes: more independent draws raise the chance of an early lucky one, which is why oversubscribing thread count still helped in Table I and why xfill’s real-grid results

are a finding-oriented result. Partitioned search instead gives genuine non-overlapping coverage, which is what a proof of unsatisfiability requires; restart-based search approximates this only through its dedicated exhaustive fallback, at the cost of forgoing restarts entirely for that one worker. This also explains why shared conflict weights helped broadly but did not resolve every grid only one architecture could solve: the mechanism improves the portfolio’s branch ordering, not what the portfolio’s fundamental coverage of the search space is.

VII. Limitations and Conclusion

All timing results are from one 14-core machine and one dictionary derived from a specific freely available word list; neither the thread-scaling numbers nor the Section V unsatisfiability result should be read as hardware- or dictionary-independent claims. Both principal solvers are actively developed software; exact commit hashes are recorded in the testbench README for reproducibility.

We described a concurrent conflict-weight-sharing mechanism for restart-portfolio CSP search occupying a specific gap in prior work – soft rather than hard information, shared continuously rather than gathered once, never used to partition – and distinguished it from the three closest points in the literature. Implemented in xfill and evaluated on a reproducible, five-solver testbench against both sophisticated competitors and unmodified naive baselines, it contributes to a broader pattern: restart portfolios and partitioned search are suited respectively to finding solutions quickly and to proving none exist, and naive backtracking without either is not competitive at real crossword scale regardless. We also reported a negative result – five heuristics rendered indistinguishable by a precondition that kept them from ever executing – as a reusable caution for comparing branching heuristics by benchmark score alone.

Acknowledgments

[Add any advisor, funding, or collaborator acknowledgments here.]

References

- [1] M. L. Ginsberg, “Dr.Fill: Crosswords and an Implemented Solver for Singly Weighted CSPs,” *Journal of Artificial Intelligence Research*, vol. 42, pp. 851–886, 2011.
- [2] C. P. Gomes, B. Selman, and H. Kautz, “Boosting Combinatorial Search Through Randomization,” in *Proc. AAAI*, 1998, pp. 431–437.
- [3] C. P. Gomes, B. Selman, N. Crato, and H. Kautz, “Heavy-tailed phenomena in satisfiability and constraint satisfaction problems,” *Journal of Automated Reasoning*, vol. 24, pp. 67–100, 2000.
- [4] F. Boussemart, F. Hemery, C. Lecoutre, and L. Sais, “Boosting systematic search by weighting constraints,” in *Proc. ECAI*, 2004, pp. 146–150.
- [5] D. Grimes and R. J. Wallace, “Sampling Strategies and Variable Selection in Weighted Degree Heuristics,” in *Proc. CP*, LNCS vol. 4741, 2007, pp. 831–838.
- [6] Y. Hamadi, S. Jabbour, and L. Sais, “ManySAT: a Parallel SAT Solver,” *J. Satisfiability, Boolean Modeling and Computation*, vol. 6, pp. 245–262, 2009.

- [7] T. Ehlers and P. J. Stuckey, “Parallelizing Constraint Programming with Learning,” in *Proc. CPAIOR*, LNCS vol. 9676, Banff, Canada, 2016.
- [8] X. Yun and S. L. Epstein, “Adaptive Parallelization for Constraint Satisfaction Search,” in *Proc. SoCS*, 2012, pp. 145–152.
- [9] X. Yun and S. L. Epstein, “A Hybrid Paradigm for Adaptive Parallel Search,” in *Proc. CP*, LNCS vol. 7514, 2012, pp. 720–736.
- [10] J.-C. Régim, M. Rezgüi, and A. Malapert, “Embarrassingly Parallel Search in Constraint Programming,” *Journal of Artificial Intelligence Research*, vol. 57, pp. 421–464, 2016.
- [11] I. P. Gent, I. J. Miguel, P. W. Nightingale, C. McCreesh, P. Prosser, N. Moore, and C. Unsworth, “A Review of Literature on Parallel Constraint Solving,” *Theory and Practice of Logic Programming*, vol. 18, no. 5–6, pp. 725–758, 2018.